

# Streamlining symbol files in Oberon

Andreas Pirklbauer

1.6.2023

## Overview

This technical note presents both a simplification and improvement of the handling of import and export for the Oberon programming language, as realized in *Extended Oberon*<sup>1</sup>, a revision of the *Project Oberon 2013* system, which is itself a reimplementation of the original *Oberon* system on an FPGA development board around 2013, as published at [www.projectoberon.com](http://www.projectoberon.com).

## Brief historical context

The topic of *symbol files* (=module interface files) has accompanied compiler development ever since the original *module* concept with *separate compilation* and type-checking *across* module boundaries (as opposed to *independent* compilation where no such checks are performed) has been introduced in the 70s and adopted in languages such as Mesa, Ada, Modula-2 and Oberon.

A correct implementation of the *module* concept was by no means obvious initially. However, the concept has evolved and today, simple implementations exist covering all key requirements, e.g.,

1. *Hidden record fields*: Offsets of non-exported pointer fields are needed for garbage collection.
2. *Re-export conditions*: Imported types may be *re-exported* and their *imports* may be hidden.
3. *Recursive data structures*: Pointer declarations may *forward reference* a record type.
4. *Module aliases*: A module can be imported under a different (alias) name.

A careful and detailed study of the evolution that led to today's status quo – which contains many useful lessons and is therefore well worth the effort – is far beyond the scope of this technical note. The reader is referred to the literature [1-13]. Here, a very rough sketch must suffice:

- Module concept introduced in 1972, early languages include Mesa, Modula and Ada [1]
- Modula-2 implementation on PDP-11 in 1979 already used *separate* compilation [2]
- Modula-2 implementation on Lilith in 1980 already used *separate* compilation [3]
- First single-pass compiler for Modula-2 compiler in 1984 used a *post-order* traversal [4, 5, 7]
- Some Oberon compilers in the 1990s used a *pre-order* traversal of the symbol table [8-11]
- The Oberon on ARM compiler (2008) used a *fixup* technique for types in symbol files [12]
- The Project Oberon 2013 compiler uses *pre-order* traversal and a *fixup* technique [13]

As with the underlying languages, all these re-implementations and refinements of the handling of import and export (and the symbol files) are characterized by a *continuous reduction of complexity*.

In this note, we present yet another simplification by eliminating the so-called “fixup” technique for *types* during export and subsequent import, as well as some additional improvements.

---

<sup>1</sup> <http://www.github.com/andreaspirklbauer/Oberon-extended>

## Symbol files in ARM Oberon (2008) and in Project Oberon (2013)

The Oberon system and compiler were re-implemented around 2013 on an FPGA development board ([www.projectoberon.com](http://www.projectoberon.com)). The compiler was derived from an earlier version of the Oberon compiler for the ARM processor. In the Project Oberon 2013 compiler, the same “fixup” technique to implement forward references *in* symbol files as in the ARM Oberon compiler is used. Quoting from the *Oberon on ARM* report [12]:

*If a type is imported again and then discarded, it is mandatory that this occurs before a reference to it is established elsewhere. This implies that types must always be defined before they are referenced. Fortunately, this requirement is fulfilled by the language and in particular by the one-pass strategy of the compiler. However, there is one exception, namely the possibility of forward referencing a record type in a pointer declaration, allowing for recursive data structures:*

```
TYPE P = POINTER TO R;  
      R = RECORD x, y: P END
```

*Hence, this case must be treated in an exceptional way, i.e. the definition of P must not cause the inclusion of the definition of R, but rather cause a forward reference in the symbol file. Such references must be fixed up when the pertinent record declaration had been read. This is the reason for the term {fix} in the syntax of (record) types. Furthermore, the recursive definition*

```
TYPE P = POINTER TO RECORD x, y: P END
```

*suggests that the test for re-import must occur before the type is established, i.e. that the type’s name must precede the type’s description in the symbol file, where the arrow marks the fixup.:*

```
TYP [#14 P form = PTR [^1]]  
TYP [#15 R form = REC [^9] lev = 0 size = 8 {y [^14] off = 4 x [^14] off = 0}] → 14
```

## Observations

The above excerpt correctly states that “types must always be defined before they are referenced”. However, if *pre-order traversal* is used when traversing the data structure called the “symbol table” to generate the symbol file – as is the case in Project Oberon 2013 – this is *already* the case.

When an identifier is to be exported, the export of the type (*Type*) precedes that of the identifier (*Object*), which therefore always refers to its type by a *backward* reference. Also, a type’s name always *precedes* the type’s description in the symbol file (see *OutType* and *Export* in *ORB*):

```
PROCEDURE OutType(VAR R: Files.Rider; t: Type);  
...  
BEGIN  
  IF t.ref > 0 THEN (*type was already output*) Write(R, -t.ref)  
  ELSE ...  
    IF t.form = Pointer THEN OutType(R, t.base)  
    ELSIF t.form = Array THEN OutType(R, t.base); ...  
    ELSIF t.form = Record THEN  
      IF t.base # NIL THEN OutType(R, t.base) ELSE OutType(R, noType) END ;  
    ELSIF t.form = Proc THEN OutType(R, t.base); ...  
  END ; ...  
END OutType;
```

```

PROCEDURE Export*(VAR modid: ORS.Ident; VAR newSF: BOOLEAN; VAR key: LONGINT);
BEGIN ...
  WHILE obj # NIL DO
    IF obj.expo THEN
      Write(R, obj.class); Files.WriteString(R, obj.name);    (*type name*)
      OutType(R, obj.type);
      IF obj.class = Typ THEN ...
      ELSIF obj.class = Const THEN ...
      END ;
      obj := obj.next
    END ;
    ...
  END Export;

```

And similarly for procedures *InType* and *Import* in *ORB*:

```

PROCEDURE InType(VAR R: Files.Rider; thismod: Object; VAR T: Type);
BEGIN Read(R, ref);
  IF ref < 0 THEN T := typtab[-ref]  (*already read*)
  ELSE NEW(t); T := t; typtab[ref] := t; t.mno := thismod.lev;
  ..
  IF form = Pointer THEN InType(R, thismod, t.base); ...
  ELSIF form = Array THEN InType(R, thismod, t.base); ...
  ELSIF form = Record THEN InType(R, thismod, t.base); ...
  ELSIF form = Proc THEN InType(R, thismod, t.base); ...
  END
END
END InType;

PROCEDURE Import*(VAR modid, modid1: ORS.Ident);
BEGIN ...
  Read(R, class);
  WHILE class # 0 DO
    NEW(obj); obj.class := class; Files.ReadString(R, obj.name);
    InType(R, thismod, obj.type); ...
    IF class = Typ THEN ...
    ELSE
      IF class = Const THEN ...
      ELSIF class = Var THEN ...
      END
    END ;
    ...
  END
END Import;

```

One can easily verify that types are *always* already “fixed” with the right value, by slightly modifying the current implementation of *ORP.Import* as follows

```

WHILE k # 0 DO
  IF typtab[k].base # t THEN ORS.Mark("type not yet fixed up") END ;
  typtab[k].base := t; Read(R, k)
END

```

The message “type not yet fixed up” will *never* be printed while importing a module. This shows that the fixup of cases of previously declared pointer types is *not necessary* as they are already “fixed” with the right value. A more formal proof can of course easily be constructed. It rests on the observation that the *type* is written to the symbol file before the corresponding *object*.

## Code that can be omitted

The following code (shown in **red**) in procedures *Import* and *Export* in ORB can be omitted.

```
PROCEDURE Import*(VAR modid, modidl: ORS.Ident);
BEGIN ...
  IF modidl = "SYSTEM" THEN
    ...
    IF F # NIL THEN
      ...
      Read(R, class);
      WHILE class # 0 DO
        NEW(obj); obj.class := class; Files.ReadString(R, obj.name);
        InType(R, thismod, obj.type); obj.lev := -thismod.lev;
        IF class = Typ THEN t := obj.type; t.typobj := obj; Read(R, k); (*<---*)
          (*fixup bases of previously declared pointer types*)
          WHILE k # 0 DO typtab[k].base := t; Read(R, k) END
        ELSE
          IF class = Const THEN ...
          ELSIF class = Var THEN ...
          END
        END
        obj.next := thismod.dsc; thismod.dsc := obj; Read(R, class)
      END
    ELSE ORS.Mark("import not available")
    END
  END
END Import;
```

```
PROCEDURE Export*(VAR modid: ORS.Ident; VAR newSF: BOOLEAN; VAR key: LONGINT);
BEGIN ...
  obj := topScope.next;
  WHILE obj # NIL DO
    IF obj.expo THEN
      Write(R, obj.class); Files.WriteString(R, obj.name);
      OutType(R, obj.type);
      IF obj.class = Typ THEN
        IF obj.type.form = Record THEN obj0 := topScope.next;
        (*check whether this is base of previously declared pointer types*)
        WHILE obj0 # obj DO
          IF (obj0.type.form = Pointer) & (obj0.type.base = obj.type)
            & (obj0.type.ref > 0) THEN Write(R, obj0.type.ref) END ;
          obj0 := obj0.next
        END
      END ;
      Write(R, 0) (*<---*)
      ELSIF obj.class = Const THEN ...
      ELSIF obj.class = Var THEN ...
      END
    END ;
    obj := obj.next
  END ;
  ...
END Export;
```

Module *ORTool* will also need to be adapted to bring it in sync with the modified module *ORB*.

## Handle alias type names among imported modules correctly

We replaced the following code in *ORB.Import*

```
obj.type.typobj := obj
```

with:

```
IF obj.type.typobj = NIL THEN obj.type.typobj := obj END
```

which establishes the *typobj* backpointer only if it doesn't exist yet. This makes sure that an imported alias type always points to the original (imported) type, not another (imported) alias type.

This is the same type of precaution as in *ORP.Declarations*, where alias types declared in the module currently being compiled are initialized as follows:

```
IF tp.typobj = NIL THEN tp.typobj := obj END
```

which also makes sure that an alias type declared in the module currently being compiled always points to the original type (no matter whether the original type is imported or declared in the module being compiled).

Without this change, compilation of module *M2* below would lead to an "undef" error at compile time when processing the declaration **VAR a: M0.TA**.

```
MODULE M0;
  TYPE TA* = RECORD i: INTEGER END ;
  TB* = TA;          (*alias type*)
END M0.

MODULE M1;
  IMPORT M0;          (*import types M0.TA and M0.TB*)
  VAR a*: M0.TA;      (*re-export type M0.TA*)
  b*: M0.TB;          (*M0.TB is not re-exported, since it is just an alias type of M0.TA*)
END M1.

MODULE M2;
  IMPORT M1, M0;      (*M1 first re-imports type M0.TA, M0 then directly imports type M0.TB*)
  VAR a: M0.TA
  b: M0.TB;
END M2.
```

The reason is that during the explicit import of module *M0* in module *M2*, the type *M0.TB* will be recognized as being the same as *M0.TA*, but the assignment *obj.type.typobj := obj* would make the backpointer for *M0.TB* point to the object node of *M0.TB* instead of leaving it as is, i.e. pointing to the original imported type *M0.TA*.

Note: The "undef" error described above can occur only if explicit imports are allowed after re-imports. In Project Oberon 2013, this is not allowed and therefore module *M2* cannot be compiled anyway ("invalid import order" error). But if it could, the "undef" error would be reported instead.

## Allow reusing the original module name if a module has been imported under an alias name

The Oberon language report defines an aliased module import as follows: *If the form "M := M1" is used in the import list, an exported object x declared within M1 is referenced in the importing module as M.x.*

Unfortunately, this definition allows several different interpretations.

### Interpretation #1:

- *It is module M1 that is imported, not M*
- *The module alias name M renames module M1*
- *A module can only have a single module alias name*

In our implementation, we have adopted this interpretation. It implies that an imported object x can be accessed only via a *single* qualified identifier *M.x* and allows reusing the original module name *M1* if a module has been imported under a module alias name *M*.

For example, the following scenarios are all legal:

```
MODULE A1; IMPORT M0 := M1, M1 := M2; END A1.
MODULE A2; IMPORT M0 := M1, M1 := M0; END A2.
MODULE A3; IMPORT M0 := M1, M2 := M0; END A3.
```

whereas the following scenario is illegal:

```
MODULE B1; IMPORT M1, A := M1, B := M1; END B1.
```

This is implemented by *not* checking the two cross-combinations *obj.orgname = name* and *obj.name = orgname* in *ORB.ThisModule*, where *obj* is an existing module in the module list.

Not checking the cross-combination *obj.orgname = name* allows the explicit import *M := M1* in:

```
MODULE A1; IMPORT M0 := M, M := M1; END A1.
```

While not checking the cross-combination *obj.name = orgname* allows the explicit imports *M := M0* and *X := M0* in:

```
MODULE A2; IMPORT M0 := M, M := M0; END A2.
MODULE A3; IMPORT M0 := M, X := M0; END A3.
```

### Alternative interpretation #2:

- *It is module M1 that is imported, not M*
- *A module alias name M is just an additional name for M1*
- *A module can have multiple module alias names*

This alternative interpretation implies that an imported object x can be accessed through *multiple* identifiers *M1.x*, *M.x*, *N.x* and does *not* allow reusing the original module name *M1* (i.e. *M1*

remains valid). This is similar to an *alias type* in Oberon – an Oberon type can have multiple alias types and the original type name remains valid.

Under this definition, modules *A1-A3* above would be illegal, whereas *B1* would be legal.

We have decided *not* to adopt interpretation #2 for the following reasons:

- Accessing the same imported object through several *different* identifiers may be confusing and is considered bad programming style. If used at all, a single module alias should suffice.
- It disallows the sometimes useful ability to "swap" two modules by writing  $M0 := M1$ ,  $M1 := M0$ .
- The already rather complex symbol table data structure in the compiler would become even more complex due to the need to now having to manage a *list* of alias objects in addition to the actual module objects. We believe this to be unnecessary complexity.

### **Allow re-imports to co-exist with module alias names and globally declared identifiers**

In the Oberon programming language, imported types can be re-exported and their original import may be hidden from the re-importing module. This means that a type from one module (M.T) can be imported by another module (M1) and then re-exported to a third module (M2) without M2 being aware of the original import from M. This can lead to hidden dependencies and create a complex hierarchy of module imports.

In Oberon, two common approaches to implement the re-export mechanism are:

1. *Self-contained symbol files*: This approach involves including imported types in symbol files, making the files self-contained and complete. This ensures that all required information is available in each symbol file, eliminating the need for recursive imports.
2. *Recursive imports*: In this approach, all required symbol files are imported recursively in full, to ensure that all necessary types are available for use. This method can result in more imports, but it is more transparent and easier to understand the dependencies between modules.

Since types that are *only* indirectly imported cannot be referred to *by name* in client modules, we have decided to *hide* re-imports from the namespace in Oberon. This allows indirectly imported types to coexist with other identifiers of the same name in the importing module. This means that:

- Module alias names of explicitly imported modules can be used without conflict.
- Globally declared identifiers of the importing module will not interfere with re-imported types.

This is implemented by simply *skipping* over re-imports in various loops that traverse the list of identifiers in the module list anchored in *ORB.topScope*.

Noting that *obj.rdo* = *TRUE* means "explicit import" and *obj.rdo* = *FALSE* means "re-import", this is achieved by adding the extra condition

```
OR (obj.class = Mod) & ~obj.rdo      (*skip over re-imports*)
```

to the guards of various loops in procedures *ORB.NewObj*, *ORB.thisObj* and *ORB.ThisModule*.

## Propagate imported export numbers of type descriptor addresses to client modules

We replaced the following code in *ORB.OutType* from the *Project Oberon 2013* implementation:

```
ELSIF t.form = Record THEN
  ...
  IF obj # NIL THEN Files.WriteNum(R, obj.exno) ELSE Write(R, 0) END ;
```

with:

```
ELSIF t.form = Record THEN
  ...
  IF obj # NIL THEN
    IF t.mno > 0 THEN Files.WriteNum(R, t.len) ELSE Files.WriteNum(R, obj.exno) END
  ELSE Write(R, 0)
  END ;
```

This makes sure that *imported* export numbers of type descriptor addresses (stored in the field *t.len*) are re-exported to client modules correctly, making it possible to perform type tests on such types when they are later re-imported.

The reader may think that this change is unnecessary, because one cannot perform a type test on re-imported types anyway (since only *explicitly* imported types can be referred to *by name* in client modules). However, our implementation *allows* explicitly importing a module *M* *after* types of *M* have previously been re-imported via other modules. In such a case, any previously re-imported type of *M* will simply be converted to an explicitly imported one in the compiler's symbol table.

Thus, we must make sure that if a type is exported by the declaring module *M* and later imported (or re-imported) and then re-exported by another module, the export number of the type descriptor address in the *declaring* module *M* is propagated through the module hierarchy correctly.

## Allow an explicit import after previous re-imports of types of the same module

In our implementation, we also allow types of a module *M* to be first re-imported and then module *M* to be explicitly imported as well. For example, in the following scenario:

- A base module *M* exports the types *T0*, *T1* and *T2*
- Module *M0* directly imports module *M* and then re-exports type *M.T0* (to clients of *M0*)
- Module *M1* directly imports module *M* and then re-exports type *M.T1* (to clients of *M1*)
- Module *C* first imports modules *M0* and *M1*, and *then* directly imports module *M*

```
MODULE C;
IMPORT
  M0,    (*re-imports type M.T0 indirectly via module M0, which itself had directly imported M*)
  M1,    (*re-imports type M.T1 indirectly via module M1, which itself had directly imported M*)
  M;     (*imports type M.T2 from M directly, since it has not yet been re-imported via other modules*)
```

the types *M.T0*, *M.T1* and *M.T2* are imported by module *C* in the following order:

- By directly importing module *M*, modules *M0* and *M1* gain access to *all* types of *M*.<sup>2</sup>
- By importing module *M0*, module *C* re-imports type *M.T0* (via module *M0*)

<sup>2</sup> Note that for the re-export mechanism to function, the symbol file of *M* must be read entirely at least once before initiating the chain of re-exports and subsequent re-imports of individual types of *M*.



- By importing module *M1*, module *C* re-imports type *M.T1* (via module *M1*)
- By importing module *M*, module *C* then imports type *M.T2* from module *M* directly

A correct implementation must make sure that for *named* types, which may be directly imported or indirectly re-imported via several symbol files (in any order), the first loaded instance (called the "primary instance") is always used and inserted into the symbol table of the compiler and the translation table *typtab* during compilation of a particular module. Furthermore, a module object for its *declaring* module must be inserted into the symbol table when a type is first encountered, regardless of whether the declaring module has already been explicitly imported or not.

If the primary instance already exists when reading a type again later from a *different* symbol file – either directly from the symbol file of the declaring module or indirectly from the symbol file of another module via the re-import mechanism – the remaining type data is read from the symbol file and then discarded.

This is necessary because we must make sure that a type which is imported or re-imported via *multiple* symbol files is always mapped to the *same* type node in the symbol table. Otherwise, incorrect incompatibilities during *type checking* may occur.

Our implementation approach is as follows:

1. Propagate the *reference number* of a type *t* (for example the types *M.T0*, *M.T1* and *M.T2*) in its *declaring* module *M* across the entire module hierarchy by means of a newly introduced field *t.orgref*. By definition, this new field has the following values:
  - If a type *t* declared in a module *M* is *directly* imported by a client module (i.e. if the client explicitly imports *M* via the `IMPORT` statement), then during compilation of the client of *M*, the field *t.orgref* is set to *t.ref*, the reference number of the type *t* in *M*. This effectively marks the start of the chain of re-exports and subsequent re-imports of the type *t*.
  - If a type *t* originally declared in module *M* is *re-imported* via an intermediate module *M0* by a client of *M0*, then during compilation of the client of *M0*, the field *t.orgref* is the reference number of the type *t* in its *declaring* module *M*, while *t.ref* is the reference number of *t* in the symbol file of module *M0*.
2. When a module *M* exports a type (for example *M.T0*) to an intermediate module *M0* and a module *C* *re-imports* this type via *M0*, a module object for its *declaring* module *M* and a type object for the imported type *M.T0* in the object list of *M*, together with its original reference number *in M*, is inserted into the symbol table of the compiler (*ORB.InType*).
3. When the same module *C* later also *explicitly* imports module *M*, we start by initializing the translation table *typtab* for module *M* with all the types of *M* that have *previously* been re-imported via other modules, using their reference numbers in the *declaring* module *M* as the index. In the above example, these are the types *M.T0* and *M.T1* (*ORB.Import*).
4. When a type declared in module *M* is read from the symbol file of *M* for the *first* time when compiling a client module, we *use* the information in the translation table to detect whether the type has previously been re-imported via *other* modules or not (*ORB.InType*).

5. If a type *has* already been re-imported from other modules *M0*, *M1*.. before, we read the type information from the symbol file of *M* and then discard it. In the above example, when module *C* imports module *M* and reads the type *M.T0* from the symbol file of *M*, it will detect that this type has already been re-imported via module *M0*. Similarly, it will detect that the type *M.T1* has already been re-imported via module *M1* (*ORB.InType*).
6. If a type has *not* been re-imported before via any other previously imported modules, we create a new entry for it in the object list of module *M*. In the above example, when module *C* imports module *M* and reads the type *M.T2* from the symbol file of *M*, we create a new entry for this type in the object list of module *M* (*ORB.InType*).
7. Before explicitly importing an object of class *Typ*, we check whether the object has already been re-imported. If so, we reuse the existing object and discard the one just read (*ORB.Import*).

### Write the module anchor before the type description to the symbol file (pre-order)

Our implementation writes the module anchor (module name and key) immediately *after* the type's reference number to the symbol file, but *before* the remaining type description, instead of *after* it as in Project Oberon 2013. This ensures that the code *also* works in the presence of *cyclic* references among types (a named type may refer to itself either directly or indirectly via type extensions).

If, for example, a record type *R1* contains a field of type *P2 = POINTER TO R2*, where the type *R2* is an extension of *R1*, then procedure *ORB.InType* will be recursively called for the types *R1* -> *P2* -> *R2* -> *R1* (in this order), where the *last* call to *ORB.InType(R, thismod, t.base)* is made when reading the type *R2*, i.e. with *t.base = R1*.

If the name and key of the module, from which the type *R1* is re-imported, were stored in the symbol file *after* the type description of *R1*, the recursion *R1* -> *P2* -> *R2* -> *R1* would also start *after* the initial assignment to *t.base* (via the assignment *T := t* at the beginning of *ORB.InType*), but *before* the code to read the name and key of the module from which *R1* is re-imported is executed. Thus, the *last* call to *InType(..., t.base)* when reading the type *R2* (i.e. with *t.base = R1*) would *also* be made *before* the code to handle the re-import of the base type *R1* is executed. But this code may *change* the field *t.base* to an entry in the translation table *typtab* (as the parameter *T* of *ORB.InType* is a *variable* parameter), thereby invalidating the initial assignment to *t.base*.

Wonders of multiple nested recursions among record and pointer base types. A little subtle indeed.

### References

1. Parnas D.L. *On the Criteria To Be Used in Decomposing Systems into Modules*. Comm ACM 15, 12 (December 1972)
2. Wirth N. *MODULA-2*. Computersysteme ETH Zürich, Technical Report No. 36 (March 1980) (ch. 15 describes the use of an implementation of Modula-2 on a DEC PDP-11 computer)
3. Geissmann L. *Separate Compilation in Modula-2 and the Structure of the Modula-2 Compiler on the Personal Computer Lilith*, ETH Zürich Dissertation No. 7286 (1983)
4. Wirth N. *A Fast and Compact Compiler for Modula-2*. Computersysteme ETH Zürich, Technical Report No. 64 (July 1985)

5. Gutknecht J. *Compilation of Data Structures: A New Approach to Efficient Modula-2 Symbol Files*. Computersysteme ETH Zürich, Technical Report No. 64 (July 1985)
6. Rechenberg, Mössenböck. *An Algorithm for the Linear Storage of Dynamic Data Structures*. Internal Paper, University of Linz (1986)
7. Gutknecht J. *Variations on the Role of Module Interfaces*. Structured Programming 10, 1, 40-46 (1989)
8. J. Templ. *Sparc-Oberon. User's Guide and Implementation*. Computersysteme ETH Zürich, Technical Report No. 133 (June 1990).
9. Griesemer R. *On the Linearization of Graphs and Writing Symbol Files*. Computersysteme ETH Zürich, Technical Report No. 156a (1991)
10. Pfister, Heeb, Templ. *Oberon Technical Notes*. Computersysteme ETH Zürich, Technical Report No. 156b (1991)
11. Franz M. *The Case for Universal Symbol Files*. Structured Programming 14: 136-147 (1993)
12. Wirth N. *An Oberon Compiler for the ARM Processor*. Technical note (December 2007, April 2008), [www.inf.ethz.ch/personal/wirth](http://www.inf.ethz.ch/personal/wirth)
13. Wirth N., Gutknecht J. *Project Oberon 2013 Edition*, [www.inf.ethz.ch/personal/wirth](http://www.inf.ethz.ch/personal/wirth)