

Streamlining symbol files in Oberon

Andreas Pirklbauer

1.8.2023

Overview

This technical note presents both a simplification and improvement of the handling of import and export for the Oberon programming language, as realized in *Extended Oberon*¹, a revision of the *Project Oberon 2013* system, which is itself a reimplementation of the original *Oberon* system on an FPGA development board around 2013, as published at www.projectoberon.com.

Brief historical context

The topic of *symbol files* (=module interface files) has accompanied compiler development ever since the original *module* concept with *separate compilation* and type-checking *across* module boundaries (as opposed to *independent* compilation where no such checks are performed) has been introduced in the 70s and adopted in languages such as Mesa, Ada, Modula-2 and Oberon.

A correct implementation of the *module* concept was by no means obvious initially. However, the concept has evolved and today, simple implementations exist covering all key requirements, e.g.,

1. *Hidden record fields*: Offsets of non-exported pointer fields are needed for garbage collection.
2. *Re-export conditions*: Imported types may be *re-exported* and their *imports* may be hidden.
3. *Recursive data structures*: Pointer declarations may *forward reference* a record type.
4. *Cyclic references*: Record and pointer types may contain cyclic references among themselves.
5. *Module aliases*: A module can be imported under a different (alias) name.

A careful and detailed study of the evolution that led to today's status quo – which contains many useful lessons and is therefore well worth the effort – is far beyond the scope of this technical note. The reader is referred to the literature [1-13]. Here, a very rough sketch must suffice:

- Module concept introduced in 1972, early languages include Mesa, Modula and Ada [1]
- Modula-2 implementation on PDP-11 in 1979 already used *separate* compilation [2]
- Modula-2 implementation on Lilith in 1980 already used *separate* compilation [3]
- First single-pass compiler for Modula-2 compiler in 1984 used a *post-order* traversal [4, 5, 7]
- Some Oberon compilers in the 1990s used a *pre-order* traversal of the symbol table [8-11]
- The Oberon on ARM compiler (2008) used a *fixup* technique for types in symbol files [12]
- The Project Oberon 2013 compiler uses *pre-order* traversal and a *fixup* technique [13]

As with the underlying languages, all these re-implementations and refinements of the handling of import and export (and the symbol files) are characterized by a *continuous reduction of complexity*.

In this note, we present yet another simplification by eliminating the so-called “fixup” technique for *types* during export and subsequent import, as well as some additional improvements.

¹ <http://www.github.com/andreaspirklbauer/Oberon-extended>

Symbol files in ARM Oberon (2008) and in Project Oberon (2013)

The Oberon system and compiler were re-implemented around 2013 on an FPGA development board (www.projectoberon.com). The compiler was derived from an earlier version of the Oberon compiler for the ARM processor. In the Project Oberon 2013 compiler, the same “fixup” technique to implement forward references *in* symbol files as in the ARM Oberon compiler is used. Quoting from the *Oberon on ARM* report [12]:

If a type is imported again and then discarded, it is mandatory that this occurs before a reference to it is established elsewhere. This implies that types must always be defined before they are referenced. Fortunately, this requirement is fulfilled by the language and in particular by the one-pass strategy of the compiler. However, there is one exception, namely the possibility of forward referencing a record type in a pointer declaration, allowing for recursive data structures:

```
TYPE P = POINTER TO R;  
R = RECORD x, y: P END
```

Hence, this case must be treated in an exceptional way, i.e. the definition of P must not cause the inclusion of the definition of R, but rather cause a forward reference in the symbol file. Such references must be fixed up when the pertinent record declaration had been read. This is the reason for the term {fix} in the syntax of (record) types. Furthermore, the recursive definition

```
TYPE P = POINTER TO RECORD x, y: P END
```

suggests that the test for re-import must occur before the type is established, i.e. that the type’s name must precede the type’s description in the symbol file, where the arrow marks the fixup.:

```
TYP [#14 P form = PTR [^1]]  
TYP [#15 R form = REC [^9] lev = 0 size = 8 {y [^14] off = 4 x [^14] off = 0}] → 14
```

Observations

The above excerpt correctly states that “types must always be defined before they are referenced”. However, if *pre-order traversal* is used when traversing the data structure called the “symbol table” to generate the symbol file – as is the case in Project Oberon 2013 – this is *already* the case.

When an identifier is to be exported, the export of the type (*Type*) precedes that of the identifier (*Object*), which therefore always refers to its type by a *backward* reference. Also, a type’s name always *precedes* the type’s description in the symbol file (see *OutType* and *Export* in *ORB*):

```
PROCEDURE OutType(VAR R: Files.Rider; t: Type);  
...  
BEGIN  
  IF t.ref > 0 THEN (*type was already output*) Write(R, -t.ref)  
  ELSE ...  
    IF t.form = Pointer THEN OutType(R, t.base)  
    ELSIF t.form = Array THEN OutType(R, t.base); ...  
    ELSIF t.form = Record THEN  
      IF t.base # NIL THEN OutType(R, t.base) ELSE OutType(R, noType) END ;  
    ELSIF t.form = Proc THEN OutType(R, t.base); ...  
  END ; ...  
END OutType;
```

```

PROCEDURE Export*(VAR modid: ORS.Ident; VAR newSF: BOOLEAN; VAR key: LONGINT);
BEGIN ...
  WHILE obj # NIL DO
    IF obj.expo THEN
      Write(R, obj.class); Files.WriteString(R, obj.name);    (*type name*)
      OutType(R, obj.type);
      IF obj.class = Typ THEN ...
      ELSIF obj.class = Const THEN ...
      END ;
      obj := obj.next
    END ;
  ...
END Export;

```

And similarly for procedures *InType* and *Import* in *ORB*:

```

PROCEDURE InType(VAR R: Files.Rider; thismod: Object; VAR T: Type);
BEGIN Read(R, ref);
  IF ref < 0 THEN T := typtab[-ref]    (*already read*)
  ELSE NEW(t); T := t; typtab[ref] := t; t.mno := thismod.lev;
  ...
  IF form = Pointer THEN InType(R, thismod, t.base); ...
  ELSIF form = Array THEN InType(R, thismod, t.base); ...
  ELSIF form = Record THEN InType(R, thismod, t.base); ...
  ELSIF form = Proc THEN InType(R, thismod, t.base); ...
  END
END
END InType;

PROCEDURE Import*(VAR modid, modid1: ORS.Ident);
BEGIN ...
  Read(R, class);
  WHILE class # 0 DO
    NEW(obj); obj.class := class; Files.ReadString(R, obj.name);
    InType(R, thismod, obj.type); ...
    IF class = Typ THEN ...
    ELSE
      IF class = Const THEN ...
      ELSIF class = Var THEN ...
      END
    END ;
  ...
END
END Import;

```

One can easily verify that types are *always* already “fixed” with the right value, by slightly modifying the current implementation of *ORP.Import* as follows

```

WHILE k # 0 DO
  IF typtab[k].base # t THEN ORS.Mark("type not yet fixed up") END ;
  typtab[k].base := t; Read(R, k)
END

```

The message “type not yet fixed up” will *never* be printed while importing a module. This shows that the fixup of cases of previously declared pointer types is *not necessary* as they are already “fixed” with the right value. A more formal proof can of course easily be constructed. It rests on the observation that the *type* is written to the symbol file before the corresponding *object*.

Code that can be omitted

The following code (shown in **red**) in procedures *Import* and *Export* in ORB can be omitted.

```
PROCEDURE Import*(VAR modid, modidl: ORS.Ident);
BEGIN ...
  IF modidl = "SYSTEM" THEN
    ...
    IF F # NIL THEN
      ...
      Read(R, class);
      WHILE class # 0 DO
        NEW(obj); obj.class := class; Files.ReadString(R, obj.name);
        InType(R, thismod, obj.type); obj.lev := -thismod.lev;
        IF class = Typ THEN t := obj.type; t.typobj := obj; Read(R, k);
          (*fixup bases of previously declared pointer types*)
          WHILE k # 0 DO typtab[k].base := t; Read(R, k) END
        ELSE
          IF class = Const THEN ...
          ELSIF class = Var THEN ...
          END
        END
        obj.next := thismod.dsc; thismod.dsc := obj; Read(R, class)
      END
    ELSE ORS.Mark("import not available")
    END
  END
END Import;

PROCEDURE Export*(VAR modid: ORS.Ident; VAR newSF: BOOLEAN; VAR key: LONGINT);
BEGIN ...
  obj := topScope.next;
  WHILE obj # NIL DO
    IF obj.expo THEN
      Write(R, obj.class); Files.WriteString(R, obj.name);
      OutType(R, obj.type);
      IF obj.class = Typ THEN
        IF obj.type.form = Record THEN obj0 := topScope.next;
          (*check whether this is base of previously declared pointer types*)
          WHILE obj0 # obj DO
            IF (obj0.type.form = Pointer) & (obj0.type.base = obj.type)
              & (obj0.type.ref > 0) THEN Write(R, obj0.type.ref) END ;
            obj0 := obj0.next
          END
        END ;
        Write(R, 0)
      ELSIF obj.class = Const THEN ...
      ELSIF obj.class = Var THEN ...
      END
    END ;
    obj := obj.next
  END ;
  ...
END Export;
```

Module *ORTool* also needs to be adapted to bring it in sync with the modified module *ORB*.

Handle alias type names among imported modules correctly

We replaced the following code in *ORB.Import*

```
obj.type.typobj := obj
```

with:

```
IF obj.type.typobj = NIL THEN obj.type.typobj := obj END
```

which establishes the *typobj* backpointer only if it doesn't exist yet. This makes sure that an imported alias type always points to the original (imported) type, not another (imported) alias type.

This is the same type of precaution as in *ORP.Declarations*, where alias types declared in the module currently being compiled are initialized as follows:

```
IF tp.typobj = NIL THEN tp.typobj := obj END
```

which also makes sure that an alias type declared in the module currently being compiled always points to the original type (no matter whether the original type is imported or declared in the module being compiled).

Without this change, compilation of module *M2* below would lead to an "undef" error at compile time when processing the declaration for variable *a*.

```
MODULE M0;
  TYPE TA* = RECORD ... END ;
  TB* = TA;                                (*alias type*)
END M0.

MODULE M1;
  IMPORT M0;                               (*import types M0.TA and M0.TB*)
  VAR a*: M0.TA;                           (*re-export type M0.TA*)
  b*: M0.TB;                               (*M0.TB is not re-exported, since it is just an alias type of M0.TA*)
END M1.

MODULE M2;
  IMPORT M1, M0;                           (*M1 first re-imports type M0.TA, M0 then directly imports type M0.TB*)
  VAR a: M0.TA;                             (*compile time error*)
  b: M0.TB;
END M2.
```

The reason is that during the explicit import of module *M0* in module *M2*, the type *M0.TB* will be recognized as being the same as *M0.TA*, but the assignment *obj.type.typobj := obj* would make the backpointer for *M0.TB* point to the object node of *M0.TB* instead of leaving it as is, i.e. pointing to the original imported type *M0.TA*.

Note: The "undef" error described above can occur only if explicit imports are allowed after re-imports. In Project Oberon 2013, this is not allowed and therefore module *M2* cannot be compiled anyway ("invalid import order" error). But if it could, the "undef" error would be reported instead.

Allow reusing the original module name if a module has been imported under an alias name

The Oberon language report defines an aliased module import as follows: *If the form "M := M1" is used in the import list, an exported object x declared within M1 is referenced in the importing module as M.x.*

Unfortunately, this definition allows several different interpretations.

Interpretation #1:

- *It is module M1 that is imported, not M*
- *The module alias name M renames module M1*
- *A module can only have a single module alias name*

In our implementation, we have adopted this interpretation. It implies that an imported object x can be accessed only via a *single* qualified identifier M.x and allows reusing the original module name M1 if a module has been imported under a module alias name M.

For example, the following scenarios are all legal:

```
MODULE A1; IMPORT M0 := M1, M1 := M2; END A1.  
MODULE A2; IMPORT M0 := M1, M1 := M0; END A2.  
MODULE A3; IMPORT M0 := M1, M2 := M0; END A3.
```

whereas the following scenario is illegal:

```
MODULE B1; IMPORT M1, A := M1, B := M1; END B1.
```

This is implemented by *not* checking the two cross-combinations *obj.orgname = name* and *obj.name = orgname* in *ORB.ThisModule*, where *obj* is an existing module in the module list.

Not checking the cross-combination *obj.orgname = name* allows the explicit import *M := M1* in:

```
MODULE A1; IMPORT M0 := M, M := M1; END A1.
```

While not checking the cross-combination *obj.name = orgname* allows the explicit imports *M := M0* and *X := M0* in:

```
MODULE A2; IMPORT M0 := M, M := M0; END A2.  
MODULE A3; IMPORT M0 := M, X := M0; END A3.
```

Alternative interpretation #2:

- *It is module M1 that is imported, not M*
- *A module alias name M is just an additional name for M1*
- *A module can have multiple module alias names*

This alternative interpretation implies that an imported object x can be accessed through *multiple* identifiers *M1.x*, *M.x*, *N.x* and does *not* allow reusing the original module name *M1* (i.e. *M1* remains valid). This is similar to an *alias type* in the Oberon language, where a type can have multiple alias types and the original type name remains valid.

Under this definition, modules *A1-A3* above would be illegal, whereas *B1* would be legal.

We have decided *not* to adopt interpretation #2 for the following reasons:

- Accessing the same imported object through several *different* identifiers may be confusing and is considered bad programming style. If used at all, a single module alias should suffice.
- It disallows the sometimes useful ability to "swap" two modules by writing *M0 := M1, M1 := M0*.
- The already rather complex symbol table data structure in the compiler would become even more complex due to the need to now having to manage a *list* of alias objects in addition to the actual module objects. We believe this to be unnecessary complexity.

Allow re-imports to co-exist with module alias names and globally declared identifiers

In the Oberon programming language, imported types can be re-exported and their original import may be hidden from the re-importing module. This means that a type from one module (*M.T*) can be imported by another module (*M1*) and then re-exported to a third module (*M2*) without *M2* being aware of the original import from *M*. This can lead to hidden dependencies and create a complex hierarchy of module imports.

In Oberon, two common approaches to implement the re-export mechanism are:

1. *Self-contained symbol files*: This approach involves including imported types in symbol files, making the files self-contained and complete. This ensures that all required information is available in each symbol file, eliminating the need for recursive imports.
2. *Recursive imports*: In this approach, all required symbol files are imported recursively in full, to ensure that all necessary types are available for use. This method can result in more imports, but it is more transparent and easier to understand the dependencies between modules.

Since types that are *only* indirectly imported cannot be referred to *by name* in client modules, we have decided to *hide* re-imports from the namespace in Oberon. This allows indirectly imported types to coexist with other identifiers of the same name in the importing module. This means that:

- Module alias names of explicitly imported modules can be used without conflict.
- Globally declared identifiers of the importing module will not interfere with re-imported types.

This is implemented by *skipping* over re-imports in various loops that traverse the list of identifiers in the module list anchored in *ORB.topScope*.

Noting that *obj.rdo = TRUE* means "explicit import" and *obj.rdo = FALSE* means "re-import", this is achieved by adding the extra condition

```
OR (obj.class = Mod) & ~obj.rdo    (*skip over re-imports*)
```

to the guards of various loops in procedures *ORB.NewObj*, *ORB.thisObj* and *ORB.ThisModule*.

Propagate imported export numbers of type descriptor addresses to client modules

We replaced the following code in *ORB.OutType* from the *Project Oberon 2013* implementation:

```
ELSIF t.form = Record THEN ...
  IF obj # NIL THEN Files.WriteNum(R, obj.exno) ELSE Write(R, 0) END ;
```

with:

```
ELSIF t.form = Record THEN ...
  IF obj # NIL THEN
    IF t.mno > 0 THEN Files.WriteNum(R, t.len) ELSE Files.WriteNum(R, obj.exno) END
  ELSE Write(R, 0)
  END ;
```

This makes sure that *imported* export numbers of type descriptor addresses (stored in the field *t.len*) are re-exported to client modules correctly, making it possible to perform type tests on such types when they are later re-imported.

The reader may think that this change is unnecessary, because one cannot perform a type test on re-imported types anyway, as only *explicitly* imported types can be referred to *by name* in clients. However, our implementation *allows* explicitly importing a module *after* types of the module have previously been re-imported via other modules. In such a case, the previously re-imported module is *converted* to an explicitly imported one in the compiler's symbol table (see *ORB.ThisModule*).

Thus, we must make sure that if a type is exported by the declaring module *M* and later imported (or re-imported) and then re-exported by another module, the export number of the type descriptor address in the *declaring* module *M* is propagated through the module hierarchy correctly.

Allow an explicit import after previous re-imports of types of the same module

In our implementation, we allow an explicit import of a module *M*, even *after* individual types of *M* have previously been re-imported via other modules. For example, in the following scenario:

```
MODULE M;
  TYPE T0* = ...; T1* = ...; T2* = ...;
END M.

MODULE M0;
  IMPORT M;                                (*import all types of M*)
  VAR t0*: M.T0;                            (*re-export type M.T0 to clients of M0*)
END M0.

MODULE M1;
  IMPORT M;                                (*import all types of M*)
  VAR t1*: M.T1;                            (*re-export M.T1 to clients of M1*)
END M1.

MODULE C;
  IMPORT
    M0,                                    (*re-import type M.T0 via module M0*)
    M1,                                    (*re-import type M.T1 via module M1*)
    M;                                     (*import type M.T2 from module M directly*)
END C.
```

the types *M.T0*, *M.T1* and *M.T2* are imported by module *C* in the following order:

- By importing module *M*, modules *M0* and *M1* gain access to *all* types of *M*.²
- By importing module *M0*, module *C* re-imports type *M.T0*
- By importing module *M1*, module *C* re-imports type *M.T1*
- By importing module *M*, module *C* imports type *M.T2* from module *M* *directly*

A correct implementation must ensure proper handling of *named* types, which may be imported or re-imported via multiple symbol files. The key requirement is to make sure that when a type is read from multiple symbol files, it is always mapped to the *same* type node in the symbol table of the compiler. Otherwise, incorrect incompatibilities during *type checking* may occur.

The basic idea of our approach is to propagate the *reference number* of a type *t* in its *declaring* module *M* across the entire module hierarchy by means of a new field *t.orgref*, defined as follows:

- If a type *t* declared in *M* is directly imported by a client module, then *t.orgref* is set to *t.ref*, i.e. the reference number of *t* in *M* itself. This effectively marks the start of the chain of re-exports and subsequent re-imports for this particular type.
- If a type *t* declared in module *M* is *re-imported* by a client via another module *M0*, then *t.orgref* is the reference number of *t* in its *declaring* module *M* (as read from the symbol file of *M0*), and *t.ref* is the reference number of *t* in the *re-exporting* module *M0*.

Thus, the reference number of type *t* in the *re-exporting* module *M0* is different from the reference number of the *same* type *t* in its *declaring* module *M*. Both are propagated to client modules.

With this preparation, our implementation approach can be described as follows:

1. When a module *M* exports a type *M.T0* and a client module *C* first *re-imports* this type via an intermediate module *M0*, a module object for its *declaring* module *M* and a type object for type *M.T0* in the object list of *M* is inserted into the symbol table of the compiler, together with the reference number of *t* in its declaring module *M* (*ORB.InType*).
2. When the same client *C* later also *explicitly* imports module *M*, we start by initializing the translation table (*typtab*) for *M* with all types of *M* that have *previously* been re-imported via other modules, using their reference numbers in the *declaring* module *M* as the index. In the above example, these are the types *M.T0* and *M.T1* (*ORB.Import*).
3. Assume module *M* is explicitly imported at a later stage. During the reading of the symbol file of *M*, when a type declaration is encountered for the first time, we *use* the information in the translation table (*typtab*) to determine whether this specific type has been previously re-imported through other modules. Importantly, since we are now importing *M* itself, we know that the reference numbers are the original ones as propagated in *t.orgref* (*ORB.InType*).
4. If a type declared in module *M* *has* already been re-imported via other modules before its direct import from *M*, we read the type information from the symbol file of *M* and then discard it. In the above example, when module *C* imports module *M*, it will detect that the type *M.T0* has already been re-imported via module *M0* and that the type *M.T1* has already been re-imported via module *M1* (*ORB.InType*).

² Note that for the re-export mechanism to function, the symbol file of *M* must be read entirely at least once before initiating the chain of re-exports and subsequent re-imports of individual types of *M*.

5. If a type declared in module *M* has *not* yet been re-imported via other modules, we create a new entry for it in the object list of *M*. This is the regular scenario, where no prior re-imports of this type have occurred. In the above example, this situation occurs when module *C* imports module *M* and reads the type *M.T2* from the symbol file of *M* (*ORB.InType*).

In summary, when an external type is encountered for the very first time during compilation of a particular module, referred to as the "primary instance" or the first loaded instance, it is inserted into the compiler's symbol table and the translation table (*typtab*). Furthermore, if the declaring module has not been explicitly imported, a corresponding module object is also created.

However, if the primary instance of a type already exists in the symbol table of the compiler, and the same type is read again from a *different* symbol file, such as the symbol file of the declaring module or another module that re-exports the type, the system handles the situation by reading the remaining type data from the symbol file currently being read and then discarding it.

Write the module anchor of re-exported type before the type description to the symbol file

When re-exporting a type, our implementation writes the module anchor (name and key) and the name of the re-exported type immediately *after* the type's reference number to the symbol file, but *before* the remaining type description instead of *after* it as in Project Oberon 2013 (*ORB.InType*).

This ensures that the code *also* works in the presence of *cyclic* references among *re-imported* types. Recall that a named type may refer to itself either directly or indirectly, either through record fields or through type *extensions*, as shown in the following example:

```
MODULE M;
  TYPE P1* = POINTER TO R1;
  P2* = POINTER TO R2;
  P3* = POINTER TO R3;
  R1* = RECORD p2*: P2 END;
  R2* = RECORD p1*: P1 END;      (*cyclic reference through record fields*)
  R3* = RECORD (R1) END;        (*cyclic reference through type extensions*)
END M.
```

Consider the case where *M* is re-imported by a module *C* indirectly via another module *M1*. In this situation, procedure *ORB.InType* is recursively called for the re-imported types *P1*, *R1*, *P2*, *R2* and *R1* (in this order) and the *last* call to *ORB.InType(R, thismod, t.base)* – made when reading the re-imported type *R2* – sets *t.base* to *R1* one recursion level up via its variable parameter *T*:

```
PROCEDURE InType(VAR R: Files.Rider; thismod: Object; VAR T: Type);
```

If the module anchor of *M* were stored in the symbol file of *M* *after* the type description of *R1* (as in Project Oberon 2013), the code to read the module anchor of *M* would also be executed *after* the recursive calls to *ORB.InType* for the re-imported types *P1*, *R1*, *P2*, *R2* and *R1*.

```
PROCEDURE InType(VAR R: Files.Rider; thismod: Object; VAR T: Type);
...
BEGIN Read(R, ref);
  IF ref < 0 THEN T := typtab[-ref]
  ELSE NEW(t); T := t; ...
    InType(R, thismod, t.base);      (*recursive calls to InType*)
    ...
    Files.ReadString(R, modname);    (*code to read module anchor*)
```

```

IF modname[0] # 0X THEN (*re-import*)
...
T := typtab[ref] (*changes the field t.base one level up in the recursion*)

```

But the code to handle re-imports may *change* the field *t.base* one level up in the recursion to an entry in the translation table (*typtab*) via the variable parameter *T* of *ORB.InType*.

Wonders of nested recursions and cyclic references among record and pointer base types! If you remain unconvinced, prepare to explore a realm of subtle intricacies ☺. Or you can choose to be on the safe side and simply make all assignments to *T* *before* any recursive calls to *ORB.InType* are initiated, as shown in the following code excerpt.

```

PROCEDURE InType(VAR R: Files.Rider; thismod: Object; VAR T: Type);
...
BEGIN Read(R, ref);
IF ref < 0 THEN T := typtab[-ref] (*already read*)
ELSE NEW(t); T := t; ...
Files.ReadString(R, modname); (*code to read module anchor*)
IF modname[0] # 0X THEN (*re-import*)
...
T := typtab[ref]
END ;
...
InType(R, thismod, t.base); (*recursive calls to InType, no more assignments to t.base*)

```

This effectively establishes the data structure for re-imported module and type objects in the symbol table *prior* to entering any recursion of *ORB.InType*. The variable parameter *T* is set to the returned type, while the local variable *t* acts as a temporary holder for data read from the currently accessed symbol file. They only differ when encountering a type that has already been read from *another* symbol file. In such cases, the subtree rooted in *t* is discarded upon re-reading the type.

References

1. Parnas D.L. *On the Criteria To Be Used in Decomposing Systems into Modules*. Comm ACM 15, 12 (December 1972)
2. Wirth N. *MODULA-2*. Computersysteme ETH Zürich, Technical Report No. 36 (March 1980) (ch. 15 describes the use of an implementation of Modula-2 on a DEC PDP-11 computer)
3. Geissmann L. *Separate Compilation in Modula-2 and the Structure of the Modula-2 Compiler on the Personal Computer Lilith*, ETH Zürich Dissertation No. 7286 (1983)
4. Wirth N. *A Fast and Compact Compiler for Modula-2*. Computersysteme ETH Zürich, Technical Report No. 64 (July 1985)
5. Gutknecht J. *Compilation of Data Structures: A New Approach to Efficient Modula-2 Symbol Files*. Computersysteme ETH Zürich, Technical Report No. 64 (July 1985)
6. Rechenberg, Mössenböck. *An Algorithm for the Linear Storage of Dynamic Data Structures*. Internal Paper, University of Linz (1986)
7. Gutknecht J. *Variations on the Role of Module Interfaces*. Structured Programming 10, 1, 40-46 (1989)
8. J. Templ. *Sparc-Oberon. User's Guide and Implementation*. Computersysteme ETH Zürich, Technical Report No. 133 (June 1990).

9. Griesemer R. *On the Linearization of Graphs and Writing Symbol Files*. Computersysteme ETH Zürich, Technical Report No. 156a (1991)
10. Pfister, Heeb, Templ. *Oberon Technical Notes*. Computersysteme ETH Zürich, Technical Report No. 156b (1991)
11. Franz M. *The Case for Universal Symbol Files*. Structured Programming 14: 136-147 (1993)
12. Wirth N. *An Oberon Compiler for the ARM Processor*. Technical note (December 2007, April 2008), www.inf.ethz.ch/personal/wirth
13. Wirth N., Gutknecht J. *Project Oberon 2013 Edition*, www.inf.ethz.ch/personal/wirth