The Extended Oberon System

Andreas Pirklbauer

1.5.2023

Extended Oberon[1] is a revision of the *Project Oberon 2013*[2] operating system and its compiler. In this document, the term *Project Oberon 2013* refers to a re-implementation of the original *Oberon* system on an FPGA development board in 2013, published at www.projectoberon.com.

Extended Oberon contains a number of enhancements, including a revision of the programming language *Oberon-2* specified as a superset of the *Oberon-07 (Revised Oberon)* language, completely *safe* unloading of modules or groups of modules and module finalization, system building and maintenance tools, smooth fractional line scrolling of displayed texts with variable line spaces, multiple logical display areas, a simple batch execution facility and various other modifications. Some of these enhancements and modifications are of a purely experimental nature, while others serve the explicit purpose of exploring potential future extensions.

1. The programming language Oberon-2 (2020 Edition)

The programming language *Oberon-2 (2020 Edition)* is a revision of the original programming language Oberon-2, as defined in 1991/1993[3]. The main difference to the original is that it is specified as a superset of *Oberon-07 (Revised Oberon)*, as defined in 2007/2016[4], rather than being based on the original language *Oberon,* as defined in 1988/1990[5].

Original Oberon (1988/1990)	➜	Original Oberon-2 (1991/1993)
↓		
Oberon-07 (2007/2016)	➜	**Oberon-2 (2020 Edition)**

Its principal new features include type-bound procedures, a dynamic heap allocation procedure for fixed-length and open arrays, a numeric and a revised type case statement, exporting and importing of string constants, no access to intermediate objects from within nested scopes and module finalization.

The language *Oberon-2 (2020 Edition)* is described in more detail in a separate document[6].

2. Safe module unloading and module finalization

Most implementations of the Oberon operating system only check for *client* references prior to module unloading. Other types of module references are usually not checked, although various approaches are typically employed to alleviate the situation where such references do exist. In

---

[1] *http://www.github.com/andreaspirklbauer/Oberon-extended*
[2] *http://www.inf.ethz.ch/personal/wirth/ProjectOberon/index.html (Project Oberon, 2013 Edition); see also http://www.projectoberon.com*
[3] *Mössenböck H., Wirth N.: The Programming Language Oberon-2. Structured Programming, 12(4):179-195, 1991*
[4] *http://www.inf.ethz.ch/personal/wirth/Oberon/Oberon07.Report.pdf (Revision 3.5.2016)*
[5] *http://inf.ethz.ch/personal/wirth/Oberon/Oberon.Report.pdf (Revision 1.10.1990)*
[6] *http://github.com/andreaspirklbauer/Oberon-extended/blob/master/Documentation/The-Revised-Oberon2-Programming-Language.pdf*

addition, the unloading of *groups* of modules with (potentially cyclic) references only among themselves and *module finalization* are usually not supported.

Checking only some, but not all references to a module prior to unloading it leaves a system in an *unsafe* state, which may – and generally will – become *unstable* the moment another module loaded later overwrites a previously released module block.

In the Oberon system, there exist several possible types of references to a loaded module M:

1. *Client references* exist when other loaded modules import module M. Client modules may refer by name to exported constants, types, variables or procedures declared in module M.

2. *Type references* exist when *type tags* (addresses of type descriptors) in dynamic objects reachable by other loaded modules refer to descriptors of types declared in module M.

3. *Procedure variable references* exist when procedure variables in static or dynamic objects reachable by other loaded modules refer to procedures declared in module M.

4. *Pointer variable references to static module data* exist when pointer variables in static or dynamic objects reachable by other loaded modules refer to *static* objects declared in module M. Such references are only possible by resorting to low-level facilities and should be avoided (pointer variables should point exclusively to anonymous variables allocated in the *heap* when needed during program execution). But since they *can* in theory exist and the required metadata (the locations of all pointer variables) already exists for other reasons, a specific implementation may opt to also check for them prior to module unloading.

Extended Oberon implements *completely safe* unloading of modules and module groups by systematically checking *all* possible types of module references from anywhere in the system.

In addition, Extended Oberon allows any module to set a *module finalization* sequence by calling procedure *Modules.Final* during *module initialization*. The specified routine must be a *parameterless* procedure declared in the *same* module. It is executed when the module is removed from the system. Alternatively, a module can use the language construct FINAL followed by the module finalization sequence. A standalone program cannot set a module finalization sequence.

The mechanism of *safe* module unloading and *module finalization* is described in more detail in a separate document[7].

3. System building and maintenance tools

A minimal version of the Oberon system *building tools*, as outlined in chapter 14 of the book *Project Oberon 2013 Edition*, has been added. They provide the necessary tools to establish the prerequisites for the *regular* Oberon startup process either on an existing *local* system or on a bare metal *target* system connected to a *host* via a data link (e.g., an RS-232 serial line).

The boot linker (procedure *ORL.Link*) links a set of object files together and generates a valid boot file from them. It can be used to either generate a *regular* boot file to be loaded onto the *boot area* of a disk or a *build-up* boot file sent over a data link to a target system. One can also

---

[7] *http://github.com/andreaspirklbauer/Oberon-extended/blob/master/Documentation/Safe-module-unloading-and-finalization-in-Oberon.pdf*

include an *entire* Oberon system in a single *boot file*. Installing it on a target is similar to booting a commercial operating system in a Plug & Play fashion over the network or from a USB stick. The command *ORL.Load* loads a valid *boot file*, as generated by the command *ORL.Link*, onto the boot area of the local disk, one of the two valid *boot sources* (the other being the data link).

The tools to build an entirely new Oberon system on a bare metal *target* system connected to a *host* system via a communication link are provided by the module pair *ORC* running on the host system and *Oberon0* running on the target. Apart from actually *building* the target system, there is a variety of other Oberon-0 commands that can be remotely initiated from the host once the Oberon-0 command interpreter is running on the target system, for example commands for system inspection, loading and unloading of modules or the remote execution of commands. Finally, there are tools to modify the boot loader itself (module *BootLoad*), a small *standalone* program permanently resident in the computer's read-only store.

The system building tools for Project Oberon 2013 and Extended Oberon are described in more detail in a separate document[8].

4. Smooth scrolling of displayed texts with variable line spaces

Extended Oberon enables completely smooth, fractional line scrolling of displayed texts with variable lines spaces and dragging of entire viewers with continuous content refresh. Both *far* (positional) scrolling and *near* (pixel-based) scrolling are realized. The system automatically switches back and forth between the two scrolling modes based on the horizontal position of the mouse pointer. For the purist, such a feature may represent a questionable embellishment of the Oberon system, but it is simply indispensable if the system is to support modern touch display devices where a mouse is absent and viewers may not have scrollbars. In such an environment, continuous scrolling is the only acceptable way to scroll and presents a more natural user interface. As a welcome side effect, the initial learning curve for users new to the Oberon system and its user interface is *considerably* reduced.

5. Multiple logical display areas ("virtual displays")

The Oberon system was designed to operate on a *single* abstract logical display area which is decomposed into a number of vertical *tracks*, each of which is further decomposed into a number of horizontal *viewers*. Extended Oberon adds the (purely experimental) ability to create *several* such display areas (or *displays* for short) on the fly and to seamlessly switch between them.

Thus, the extended conceptual hierarchy of the display system consists of the triplet *(display, track, viewer)* and consequently the underlying base module *Viewers* exports procedures to add and remove *displays*, open and close *tracks* within existing displays and open and close individual *viewers* within tracks. There are no restrictions on the number of displays, tracks or viewers that can be created. Focus viewers and text selections are separately defined for each display, but there is only one system-wide log.

This scheme naturally maps to systems with multiple *physical* monitors attached to the same computer. It can also be used to realize super-fast context switching, for example in response to a swipe gesture on a touch display device.

---

[8] http://github.com/andreaspirklbauer/Oberon-extended/blob/master/Documentation/The-Oberon-system-building-tools.pdf

The command *System.OpenDisplay name* opens a new logical display with the specified name, *System.CloseDisplay id* closes an existing one. *System.ShowDisplays* lists all open displays, *System.ThisDisplay* shows the display *id* and *name* of the current display, *System.SetDisplay id* switches to a new display, *System.SetDisplayName id name* assigns a new name to an existing display, and *System.PrevDisplay* and *System.NextDisplay* "rotate" through the open displays.

The additional commands *System.Expand, System.Spread* and *System.Clone* are displayed in the title bar of every menu viewer. *System.Expand* expands the viewer *as much as possible* by reducing *all* other viewers in its track to their minimum heights, leaving just their title bars visible. The user can switch back to any of the minimized viewers by clicking on *System.Expand* again in any of these title bars. *System.Spread* evenly redistributes all viewers vertically. This may be useful after having invoked *System.Expand*. *System.Clone* opens a new logical display on the fly *and* displays a copy of the initiating viewer *there*. The user can then toggle between the two copies of the viewer (i.e., switch logical *displays*) with a single mouse click.

6. Simple batch execution facility

The command *System.Batch file* executes multiple commands separated by ~ in the specified batch file. The contents of the batch file are displayed in a system viewer. Execution stops when the first invalid command is encountered (this can be another ~). Comments are not allowed.

For example, if a batch file *BuildInnerCore.Batch* is created containing the commands

```
ORP.Compile Kernel.Mod/s Disk.Mod/s FileDir.Mod/s Files.Mod/s Modules.Mod/s ~
ORL.Link Modules ~
~
```

then the command *System.Batch BuildInnerCore.Batch ~* will build the Oberon *inner core*.

The command *Oberon.Batch* is almost identical to *System.Batch*, except that it executes the commands *immediately following* it in the text, where the command itself is located *and* that it invokes the garbage collector after *each* executed command, if necessary[9].

The output of the command *Oberon.Batch* is displayed in the system-wide log. For example, the following command in the tool file *System.Tool* allows the user to build the entire Oberon system with a *single* mouse click:

```
Oberon.Batch
   ORP.Compile Kernel.Mod Disk.Mod FileDir.Mod Files.Mod Modules.Mod ~
   ORL.Link Modules ~
   ORL.Load Modules.bin ~
   ORP.Compile Input.Mod/s Display.Mod/s Viewers.Mod/s ~
   ORP.Compile Fonts.Mod/s Texts.Mod/s ~
   ORP.Compile Oberon.Mod/s ~
   ORP.Compile MenuViewers.Mod/s ~
   ORP.Compile TextFrames.Mod/s ~
   ORP.Compile System.Mod/s Edit.Mod/s ~
   ORP.Compile Tools.Mod/s FontTool.Mod/s ~
   ORP.Compile ORS.Mod/s ORB.Mod/s ~
   ORP.Compile ORG.Mod/s ORP.Mod/s ~
   ORP.Compile ORL.Mod/s ORX.Mod/s ORTool.Mod/s ~
 ~
```

---

[9] *The reason why System.Batch does not invoke the garbage collector is that it allocates a text and a system viewer (referred to by local variables) which would be collected. We could have chosen to make these variables global, but we opted to restrict automatic garbage collection to the command Oberon.Batch (where the stack is known to not contain pointers into the heap).*

Finally, Extended Oberon executes the batch file *System.Batch* when the system starts, e.g.,

```
PCLink1.Run ~
System.Watch ~
~
```

7. Various other modifications

Various other modifications have been made to the Oberon compiler and runtime system, some of which are of a purely experimental nature. For example, in the original *Project Oberon 2013* operating system, the following RISC code computes the address of an *external* variable:

```
LDR RH, MT, mno*4              base address of data section of the imported module
ADD RH, RH, offset             offset computed by the loader from object's export number
```

where the base address is fetched from a table global to the entire system. This table contains one entry for every module loaded, namely the address of the module's data section called the module's *static base*. The address of this table is permanently held in register MT (=R12). We note that this instruction sequence imposes a restriction of 64KB for the offset from the static base of an *imported* module, as the ADD instruction has an operand field of 16 bits only.

The same two-instruction sequence is used to compute the address of a global variable of the *same* module (mno=0), but only if the offset from the module's static base is *less* than 64KB.

If this offset is *greater* than 64KB, the following four-instruction sequence is generated instead:

```
LDR RH, MT, mno*4                    base address of the module's data section (static base)
MOV' RH+1, offset DIV 10000H         RH+1[16..31] := upper 16 bits of the offset from the static base
IOR RH+1, RH+1, offset MOD 10000H    RH+1[0..15] := lower 16 bits of the offset from the static base
ADD RH, RH, RH+1                     RH := static base + offset from the static base
```

Extended Oberon uses neither a global module table nor a register to hold the base address of the data section of a module at run time. Instead, *all* global and external data references are encoded in a special format in the object file, and are replaced by *absolute* addresses at module load time[10], always resulting in the following two-instruction sequence for computing the address of a global or external variable:

```
MOV' RH, (staticbase + offset) DIV 10000H      RH[16..31] := upper 16 bits of (staticbase + offset)
IOR RH, RH, (staticbase + offset) MOD 10000H   RH[0..15] := lower 16 bits of (staticbase + offset)
```

We emphasize that this instruction sequence does not access any memory at all. If the value of the variable is to be loaded into a register, the IOR instruction is replaced by an LDR instruction.

This approach has several advantages:

First, accessing a global variable of the *same* module is no different than accessing an *external* variable, as the *same* two-instruction sequence shown above is generated in both cases. When accessing a global variable of the *same* module at an offset *greater* than 64KB from the base address of its data section, this results in higher code density and faster code execution.

---

[10] If a data element of the s a m e module is accessed and the offset from the base address of the module's data section is greater than 64KB (requiring more than 16 bits), the offset is initially spread across the two-instruction sequence during compilation (8 bits in the first, 16 bits in the second) and reassembled – and converted to an absolute address – by the loader.

Second, there is no more 64KB restriction on the offset from the base address of the data section of *imported* modules, as the ADD instruction (with a 16-bit offset field) is no longer used.

Third, it eliminates the need for a *module table* global to the entire system and with it also the need for a dedicated MT (=R12) register permanently holding the address of this table.

Finally, it eliminates the need for a static base *register* holding the address of a module's data section at run time. In earlier versions of the *Project Oberon 2013* system, this base address was held in a dedicated register SB (=R13). In later versions, it was replaced by an arbitrary register. In Extended Oberon, the base address of a module's data section is only used during loading and linking, but *not* at execution time, as the module loader always inserts *absolute* addresses for *all* global and external data references.

* * *